

EECS 498: Reinforcement Learning

Homework 5 Responses

Tejas Jha
tjha

December 12, 2018

This document includes my responses to Homework 5 questions. Responses that involved the use of coding will provide references to specific lines of code to provide a better overview of how the problem was approached. The code can either be referenced in the Appendix or in the accompanied python script submitted with this assignment.

Question 1

- (a)
- (b)

Suppose the reward function for an MDP is a linear function of d features for a state s

$$R(s) = \alpha_1\phi_1(s) + \alpha_2\phi_2(s) + \dots + \alpha_d\phi_d(s)$$

where the $\phi_1 \dots \phi_d$ are fixed, known and bounded basis functions mapping from state space S to the reals.

As presented in the Algorithms for Inverse Reinforcement Learning paper, we can define a value function for a policy π that maps $S \mapsto A$ for any state s_1 as the following:

$$V^\pi(s_1) = E[R(s_1) + \gamma R(s_2) + \gamma^2 R(s_3) + \dots | \pi]$$

where the expectation is over the distribution of the state sequence (s_1, s_2, \dots) .

Additionally from the same paper, we can use the notation V_i^π to denote the value function of the policy π in the MDP when the reward function is $R = \phi_i$. We can prove that for any policy π , the value function can be defined as $V^\pi(s_1) = \alpha_1 V_1^\pi + \alpha_2 V_2^\pi + \dots + \alpha_d V_d^\pi$

We can substitute the reward function into the expectation in the value function equation and use the linearity of expectation to expand it to a sum of expectations of fixed and known values. These expectations would simplify to just the terms. This is shown below:

$$\begin{aligned} V^\pi(s_1) &= E[R(s_1) + \gamma R(s_2) + \gamma^2 R(s_3) + \dots | \pi] \\ V^\pi(s_1) &= E[\alpha_1\phi_1(s_1) + \alpha_2\phi_2(s_1) + \dots + \alpha_d\phi_d(s_1)] \end{aligned}$$

$$+ \gamma * (\alpha_2 \phi_1(s_2) + \alpha_2 \phi_2(s_2) + \dots + \alpha_d \phi_d(s_2)) + \\ \gamma^2 * (\alpha_1 \phi_1(s_3) + \alpha_2 \phi_2(s_3) + \dots + \alpha_d \phi_d(s_3)) + \dots | \pi]$$

$$V^\pi(s_1) = \alpha_1 \phi_1(s_1) + \alpha_2 \phi_2(s_1) + \dots + \alpha_d \phi_d(s_1) \\ + \gamma * (\alpha_2 \phi_1(s_2) + \alpha_2 \phi_2(s_2) + \dots + \alpha_d \phi_d(s_2)) + \\ \gamma^2 * (\alpha_1 \phi_1(s_3) + \alpha_2 \phi_2(s_3) + \dots + \alpha_d \phi_d(s_3)) + \dots$$

The terms can now be regrouped using the alpha terms.

$$V^\pi(s_1) = \alpha_1 * (\phi_1(s_1) + \gamma \phi_1(s_2) + \gamma^2 \phi_1(s_3) + \dots) \\ + \alpha_2 * (\phi_2(s_1) + \gamma \phi_2(s_2) + \gamma^2 \phi_2(s_3) + \dots) \\ + \dots \\ + \alpha_d * (\phi_d(s_1) + \gamma \phi_d(s_2) + \gamma \phi_d(s_3) + \dots)$$

$$V^\pi(s_1) = \alpha_1 * E[\phi_1(s_1) + \gamma \phi_1(s_2) + \gamma^2 \phi_1(s_3) + \dots | \pi] \\ + \alpha_2 * E[\phi_2(s_1) + \gamma \phi_2(s_2) + \gamma^2 \phi_2(s_3) + \dots | \pi] \\ + \dots \\ + \alpha_d * E[\phi_d(s_1) + \gamma \phi_d(s_2) + \gamma^2 \phi_d(s_3) + \dots | \pi]$$

Using the definition of V_i^π and the equation for the value function, we can simplify the above expression to:

$$V^\pi(s_1) = \alpha_1 V_1^\pi + \alpha_2 V_2^\pi + \dots + \alpha_d V_d^\pi$$

Since s_1 is a variable to represent any arbitrary state, we can also just write the above as:

$$V^\pi(s) = \alpha_1 V_1^\pi + \alpha_2 V_2^\pi + \dots + \alpha_d V_d^\pi$$

This concludes the proof.

Question 2

Appendix: Relevant Code - tjha_hw5.py

```
1 # Tejas Jha
2 # EECS 498 — Reinforcement Learning
3 # Homework 5
4 # 12 December 2018
5
6
7 import gym
```

```

8 from gym import wrappers
9 import random
10 import math
11
12 from keras.models import Sequential
13 from keras.layers import Dense, Activation
14 from keras.optimizers import Adam
15 import matplotlib.pyplot as plt
16
17 import numpy as np
18
19 # hyper parameters
20 EPISODES = 100 # number of episodes
21 EPS_START = 0.5 # e-greedy threshold start value
22 EPS_END = 0.01 # e-greedy threshold end value
23 EPS_DECAY = 0.0001 # e-greedy threshold decay
24 GAMMA = 1.0 # Q-learning discount factor
25 LR = 0.0001 # NN optimizer learning rate
26 HIDDEN_LAYER = 64 # NN hidden layer size
27 BATCH_SIZE = 32 # Q-learning batch size
28
29
30 class ReplayMemory:
31     def __init__(self, capacity):
32         self.capacity = capacity
33         self.memory = []
34
35     def push(self, transition):
36         self.memory.append(transition)
37         if len(self.memory) > self.capacity:
38             del self.memory[0]
39
40     def sample(self, batch_size):
41         return random.sample(self.memory, batch_size)
42
43     def __len__(self):
44         return len(self.memory)
45
46
47 class Network(nn.Module):
48     def __init__(self):
49         nn.Module.__init__(self)
50         self.l1 = nn.Linear(4, HIDDEN_LAYER)
51         self.l2 = nn.Linear(HIDDEN_LAYER, 2)
52
53     def forward(self, x):
54         x = F.relu(self.l1(x))
55         x = self.l2(x)

```

```

56         return x
57
58
59 env = gym.make('CartPole-v0')
60 env = wrappers.Monitor(env, './tmp/cartpole-v0-1')
61
62 model = Network()
63 if use_cuda:
64     model.cuda()
65 memory = ReplayMemory(10000)
66 optimizer = optim.Adam(model.parameters(), LR)
67 steps_done = 0
68 episode_durations = []
69
70
71 def select_action(state):
72     global steps_done
73     sample = random.random()
74     eps_threshold = EPS_END + (EPS_START - EPS_END) * math.exp(-1. *
75         steps_done / EPS_DECAY)
76     if sample > eps_threshold:
77         return model.Variable(state, volatile=True).type(FloatTensor).
78             data.max(1)[1].view(1, 1)
79     else:
80         return LongTensor([random.randrange(2)])
81
82 def run_episode(e, environment):
83     state = environment.reset()
84     steps = 0
85     while True:
86         environment.render()
87         action = select_action(FloatTensor([state]))
88         next_state, reward, done, _ = environment.step(action[0, 0])
89
90         # negative reward when attempt ends
91         if done:
92             reward = -1
93
94         memory.push((FloatTensor([state]),
95             action, # action is already a tensor
96             FloatTensor([next_state]),
97             FloatTensor([reward])))
98
99         learn()
100
101     state = next_state

```

```

102         steps += 1
103
104         if done:
105             print("{2}_Episode_{0}_finished_after_{1}_steps"
106                   .format(e, steps, '\033[92m' if steps >= 195 else '
107                               \033[99m'))
108             episode_durations.append(steps)
109             plot_durations()
110             break
111
112 def learn():
113     if len(memory) < BATCH_SIZE:
114         return
115
116     # random transition batch is taken from experience replay memory
117     transitions = memory.sample(BATCH_SIZE)
118     batch_state, batch_action, batch_next_state, batch_reward = zip(*
119                               transitions)
120
121     batch_state = Variable(torch.cat(batch_state))
122     batch_action = Variable(torch.cat(batch_action))
123     batch_reward = Variable(torch.cat(batch_reward))
124     batch_next_state = Variable(torch.cat(batch_next_state))
125
126     # current Q values are estimated by NN for all actions
127     current_q_values = model(batch_state).gather(1, batch_action)
128     # expected Q values are estimated from actions which gives maximum
129     # Q value
130     max_next_q_values = model(batch_next_state).detach().max(1)[0]
131     expected_q_values = batch_reward + (GAMMA * max_next_q_values)
132
133     # loss is measured from error between current and newly expected Q
134     # values
135     loss = F.smooth_l1_loss(current_q_values, expected_q_values)
136
137     # backpropagation of loss to NN
138     optimizer.zero_grad()
139     loss.backward()
140     optimizer.step()
141
142 def plot_durations():
143     plt.figure(2)
144     plt.clf()
145     durations_t = torch.FloatTensor(episode_durations)
146     plt.title('Training ...')
147     plt.xlabel('Episode')

```

```

146     plt.ylabel('Duration')
147     plt.plot(durations_t.numpy())
148     # take 100 episode averages and plot them too
149     if len(durations_t) >= 100:
150         means = durations_t.unfold(0, 100, 1).mean(1).view(-1)
151         means = torch.cat((torch.zeros(99), means))
152         plt.plot(means.numpy())
153
154     plt.pause(0.001) # pause a bit so that plots are updated
155
156
157 for e in range(EPIISODES):
158     run_episode(e, env)
159
160 print('Complete')
161 env.render(close=True)
162 env.close()
163 plt.ioff()
164 plt.show()

```